

# COL380 Assignment 4

## GPU Matrix Multiplication

Nischay Diwan  
2020CS50433

April 2023

### 1 Approach 1

The initial approach of the algorithm was using the compressed sparse row matrix representation of Matrix A. Then the algorithm iterates over all the blocks of the resultant matrix initialized to zero and computes this block. Here for each block of the resultant matrix, we iterate over only non-zero blocks of Matrix A and then find the corresponding block if it exists in Matrix B. Then, if the blocks are non-zero, it is sent to the device for parallel multiplication for this block. So, only 1 block is computed simultaneously but by  $m^2$  parallel threads. If the block is non-zero, the output is written. This algorithm performs equivalently or slower than using a CPU.

### 2 Approach 2

First, the cuda memory is allocated repeatedly and transfers only 2 blocks of data. This is changed to allocating and sending a complete matrix(non-zero blocks only) once at the start. This finally brings some speedup into the algorithm. But the main change in approach 2 is to allow parallel computation of all the blocks. So each block in  $n^2/m^2$  output is computed by each block of GPU using  $m^2$  threads. This brings about 25-30x speedup. It also allows scalability on large sizes and dense matrices.

### 3 Approach 3

Final cuda memory allocation and transfer code:

```
1 void matMul(vector<array<int,3>> &mp1, vector<uint> &blksA, vector<array<int,3>> &mp2,  
    vector<uint> &blksB, long long n, long long m, vector<uint> &blksC){  
2     long long nm = n/m;  
3     // sending data to GPU  
4     int streamSize = 2;  
5     cudaError_t err;  
6     cudaStream_t stream[streamSize];  
7     for(int i = 0;i<streamSize;i++){  
8         cudaStreamCreate(&stream[i]);  
9     }  
10    size_t size = sizeof(uint);  
11    size_t size2 = sizeof(int);  
12    size_t size3 = sizeof(uint) * n * n;  
13    uint *a = &blksA[0], *b = &blksB[0], *da, *db;  
14    uint *c = &blksC[0], *dc;  
15    cudaMalloc(&da,size*blksA.size());  
16    cudaMalloc(&db,size*blksB.size());  
17    cudaMalloc(&dc,size3);  
18    cudaDeviceSynchronize();  
19    cudaMemset(dc,0,size3);  
20    cudaMemcpyAsync(da,a,(size_t)size*(size_t)blksA.size(),cudaMemcpyHostToDevice,stream  
    [0]);  
21    cudaMemcpyAsync(db,b,(size_t)size*(size_t)blksB.size(),cudaMemcpyHostToDevice,stream  
    [1]);
```

Initially, 2 cuda streams are generated, and both matrices in non-zero blocks are transferred, as shown in the above code. Here initial memory allocation is done. 2 streams help in parallel data transfers.

```

22 // converting to CSR
23 vector<int> valV;
24 vector<int> colV;
25 vector<int> rofV;
26 int offset = 0;
27 int rowno = 0;
28 rofV.push_back(offset);
29 for(int i = 0; i < mp1.size() ; i++){
30     colV.push_back(mp1[i][1]);
31     valV.push_back(mp1[i][2]);
32     if(mp1[i][0] > rowno){
33         for(int cc = 0; cc < (mp1[i][0] - rowno); cc++){
34             rofV.push_back(offset);
35         }
36         rowno = mp1[i][0];
37     }
38     offset+=1;
39 }
40 for(int j = rowno; j < nm; j++){
41     rofV.push_back(blksA.size()/m/m);
42 }
43 vector<int> valV2;
44 vector<int> cofV;
45 vector<int> rowV;
46 offset = 0;
47 int colno = 0;
48 cofV.push_back(offset);
49 for(int i = 0; i < mp2.size() ; i++){
50     rowV.push_back(mp2[i][0]);
51     valV2.push_back(mp2[i][2]);
52     if(mp2[i][1] > colno){
53         for(int cc = 0; cc < (mp2[i][1] - colno); cc++){
54             cofV.push_back(offset);
55         }
56         colno = mp2[i][1];
57     }
58     offset+=1;
59 }
60 for(int j = colno; j < nm; j++){
61     cofV.push_back(blksB.size()/m/m);
62 }
63 std::cout << "CSR converted\n";

```

In the above part of the code, the non-zero blocks are iterated, and matrices A and B are converted to Compressed Sparse Matrix Representations. The difference here is that both are converted to sparse form, and later, binary search is replaced with 2 pointer approach (this is a major speedup: about 10-15x). This also reduces the data to be transferred and also faster.

```

64 int *ka;
65 cudaMalloc(&ka, (size_t) size2 * (size_t) mp1.size());
66 cudaMemcpyAsync(ka, valV.data(), (size_t) size2 * (size_t) mp1.size(), cudaMemcpyHostToDevice,
67     stream[0]);
68 int *kb;
69 cudaMalloc(&kb, (size_t) size2 * (size_t) mp2.size());
70 cudaMemcpyAsync(kb, valV2.data(), (size_t) size2 * (size_t) mp2.size(), cudaMemcpyHostToDevice,
71     stream[1]);
72 int *rof, *col, *cof, *row;
73 cudaMalloc(&rof, (size_t) rofV.size() * size2);
74 cudaMalloc(&col, (size_t) colV.size() * size2);
75 cudaMalloc(&cof, (size_t) cofV.size() * size2);
76 cudaMalloc(&row, (size_t) rowV.size() * size2);
77 cudaMemcpyAsync(rof, &rofV[0], (size_t) rofV.size() * size2, cudaMemcpyHostToDevice, stream
78     [0]);
79 cudaMemcpyAsync(col, &colV[0], (size_t) colV.size() * size2, cudaMemcpyHostToDevice, stream
80     [1]);
81 cudaMemcpyAsync(cof, &cofV[0], (size_t) cofV.size() * size2, cudaMemcpyHostToDevice, stream

```

```

[0]);
78  cudaMemcpyAsync(row,&rowV[0],(size_t)rowV.size()*size2,cudaMemcpyHostToDevice,stream
[1]);

```

After this remaining CSMR data is sent asynchronously, the device kernel is invoked with  $2 * m^2 * \text{sizeof}(\text{uint})$  of shared memory for each block,  $n^2/m^2$  blocks and  $m^2$  threads per block. After this, the cudamemory is freed.

```

79  int stride = (int)(nm*nm);
80  cudaDeviceSynchronize();
81  matMulGPU<<<stride,m*m,2*size*m*m,0>>>(da,db,dc,ka,kb,rof,col,row,cof,m,n,mp1.size(),
    mp2.size()); // i X k
82  cudaMemcpy((void *)c,(void *)dc,size3,cudaMemcpyDeviceToHost);
83  // free the memory
84  cudaFree(da);
85  cudaFree(db);
86  cudaFree(ka);
87  cudaFree(kb);
88  cudaFree(rof);
89  cudaFree(col);
90  cudaFree(dc);
91  for(int i = 0;i<streamSize ;i++){
92      cudaStreamDestroy(stream[i]);
93  }
94  }

```

Final device kernel code:

```

1  __global__
2  void matMulGPU(uint *a, uint *b, uint *c, int *ka, int *kb, int *rof, int *col, int m,
    int n, int k1, int k2){
3      extern __shared__ uint dab[];
4      int bid = blockIdx.x;
5      int tid = threadIdx.x;
6      int nm = n/m;
7      int i = bid / nm;
8      int k = bid % nm;
9      uint64_t temp = 0;
10     for (int j = rof[i]; j < rof[i+1]; j++){
11         int id1 = ka[j];
12         int cl = col[j];
13         int id2 = binASearch(kb,k2,cl,k);
14         if(!(id2 == -1)){
15             dab[tid] = (uint)a[tid + id1*m*m];
16             dab[tid + m*m] = (uint)b[tid + id2*m*m];
17             __syncthreads();
18             int ii = tid/m;
19             int jj = tid%m;
20             for (int kk = 0; kk < m; ++kk)
21             {
22                 temp = temp + (uint64_t)(dab[ii*m + kk] * dab[kk*m + jj + m*m]);
23             }
24             __syncthreads();
25         }
26         __syncthreads();
27     }
28     __syncthreads();
29     c[tid + i*m*n + k*m*m] = min(temp,MAX_VAL);
30 }

```

Here simply, each block bid in  $n^2/m^2$  blocks multiplies the corresponding blocks of A and B to make the resultant block. The  $m^2$  threads individually calculate each block element within each block multiplication.

## 4 Notes and Other Optimizations

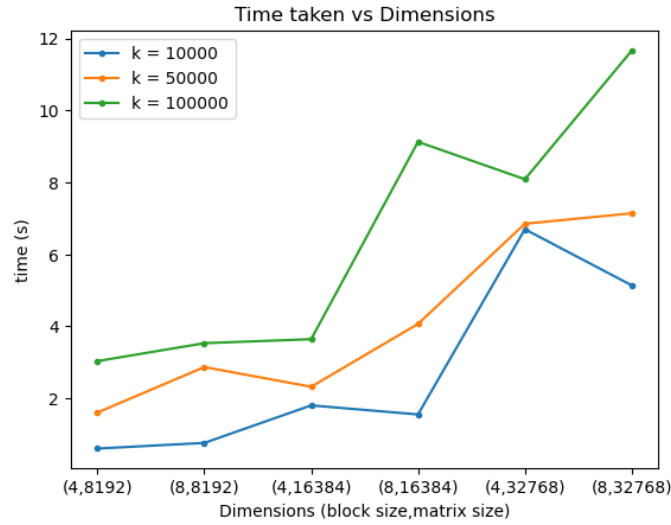
- Earlier stl maps were used to read and store the input but are now replaced with vectors of int3 and unit that store the index of the form block's row, column and array index and a vector that stores  $k1 * m^2$  non zero block values.

- Experimented with grid dimensions and block sizes; however, that did not bring any significant speedup. So, the GPU kernel uses  $2 * m^2 * \text{sizeof}(uint)$  shared memory for each block; it uses  $n^2/m^2$  blocks and  $m^2$  threads per block.

## 5 Analysis

### 5.1 Time for each test-cases

(Nodes,Threads)	(4,8192)	(8,8192)	(4,16384)	(8,16384)	(4,32768)	(8,32768)
Test1	0.605	0.762	1.804	1.553	6.697	5.139
Test2	1.606	2.869	2.322	4.067	6.850	7.139
Test3	3.034	3.533	3.643	9.125	8.086	11.659



### 5.2 Large Tests

Input reading done, n: 32768, m: 8, k-values: 8800000 and 88000000

Sorted input: 50.2979

==1170== NVPROF is profiling process 1170, command: ./exec input1 input2 outFile.bin

CSR converted

Time taken for matrix multiplication: 73.9428

Writing output

Number of output non-zero blocks: 16777216

Writing done

==1170== Profiling application: ./exec input1 input2 outFile.bin

==1170== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	97.38%	71.2514s	1	71.2514s	71.2514s	71.2514s	matMulGPU(unsigned int*, unsigned int*, unsigned int*, int*, int*, int*, int*, int*, int*, int, int, int, int)
	1.35%	986.72ms	8	123.34ms	3.4560us	479.70ms	[CUDA memcpy HtoD]
	1.27%	926.46ms	1	926.46ms	926.46ms	926.46ms	[CUDA memcpy DtoH]
	0.01%	7.3874ms	1	7.3874ms	7.3874ms	7.3874ms	[CUDA memset]
API calls:	98.18%	72.1783s	1	72.1783s	72.1783s	72.1783s	cudaMemcpy
	1.35%	995.15ms	8	124.39ms	20.320us	485.51ms	cudaMemcpyAsync
	0.41%	298.98ms	2	149.49ms	3.3510us	298.97ms	cudaStreamCreate
	0.04%	31.351ms	7	4.4788ms	7.3040us	10.615ms	cudaFree
	0.01%	10.379ms	9	1.1533ms	4.2550us	4.3703ms	cudaMalloc
	0.00%	381.67us	1	381.67us	381.67us	381.67us	cuDeviceTotalMem
	0.00%	164.09us	101	1.6240us	137ns	69.008us	cuDeviceGetAttribute
	0.00%	56.581us	2	28.290us	8.6290us	47.952us	cudaDeviceSynchronize
	0.00%	42.429us	1	42.429us	42.429us	42.429us	cudaLaunchKernel
	0.00%	38.825us	1	38.825us	38.825us	38.825us	cudaMemset
	0.00%	29.269us	1	29.269us	29.269us	29.269us	cuDeviceGetName
	0.00%	23.936us	2	11.968us	3.4970us	20.439us	cudaStreamDestroy
	0.00%	9.3970us	1	9.3970us	9.3970us	9.3970us	cuDeviceGetPCIBusId
	0.00%	1.6910us	3	563ns	274ns	1.1100us	cuDeviceGetCount
	0.00%	1.3800us	2	690ns	226ns	1.1540us	cuDeviceGet
	0.00%	275ns	1	275ns	275ns	275ns	cuDeviceGetUuid

```

real 2m50.428s
user 2m31.280s
sys 0m15.662s

```

```

Input reading done, n: 32768, m: 4, k-values: 1000000 and 1000000
Sorted input: 2.02364
==386== NVPROF is profiling process 386, command: ./exec input1 input2 outFile.bin
CSR converted
Time taken for matrix multiplication: 6.78835
Writing output
Number of output non-zero blocks: 56225712
Writing done
==386== Profiling application: ./exec input1 input2 outFile.bin
==386== Profiling result:

```

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:		84.08%	5.33634s	1	5.33634s	5.33634s	5.33634s	matMulGPU(unsigned int*, unsigned int*, unsigned int*, int*, int*, int*, int*, int, int, int, int)
		15.16%	961.98ms	1	961.98ms	961.98ms	961.98ms	[CUDA memcpy DtoH]
		0.47%	29.659ms	8	3.7073ms	4.9920us	13.585ms	[CUDA memcpy HtoD]
		0.29%	18.442ms	1	18.442ms	18.442ms	18.442ms	[CUDA memset]
API calls:		95.86%	6.29870s	1	6.29870s	6.29870s	6.29870s	cudaMemcpy
		3.15%	207.14ms	2	103.57ms	3.6860us	207.13ms	cudaStreamCreate
		0.74%	48.943ms	8	6.1179ms	21.123us	31.733ms	cudaMemcpyAsync
		0.12%	8.1196ms	7	1.1599ms	6.3680us	4.5655ms	cudaFree
		0.10%	6.4531ms	9	717.01us	6.8990us	4.9123ms	cudaMalloc
		0.01%	711.94us	2	355.97us	351.97us	359.98us	cuDeviceTotalMem
		0.01%	601.26us	202	2.9760us	129ns	133.62us	cuDeviceGetAttribute
		0.00%	83.458us	2	41.729us	9.1930us	74.265us	cudaDeviceSynchronize
		0.00%	72.175us	2	36.087us	24.986us	47.189us	cuDeviceGetName
		0.00%	30.094us	1	30.094us	30.094us	30.094us	cudaMemset
		0.00%	28.433us	1	28.433us	28.433us	28.433us	cudaLaunchKernel
		0.00%	15.683us	2	7.8410us	3.3400us	12.343us	cudaStreamDestroy
		0.00%	8.1140us	2	4.0570us	1.3880us	6.7260us	cuDeviceGetPCIBusId
		0.00%	2.2730us	4	568ns	165ns	1.4900us	cuDeviceGet
		0.00%	2.1840us	3	728ns	301ns	1.5370us	cuDeviceGetCount
		0.00%	493ns	2	246ns	229ns	264ns	cuDeviceGetUuid

```

real 0m49.730s
user 0m41.582s
sys 0m7.030s

```

```

Input reading done, n: 32768, m: 4, k-values: 10000000 and 10000000
Sorted input: 21.1806
==222== NVPROF is profiling process 222, command: ./exec input1 input2 outFile.bin
CSR converted
Time taken for matrix multiplication: 66.5818
Writing output
Number of output non-zero blocks: 67108864
Writing done
==222== Profiling application: ./exec input1 input2 outFile.bin
==222== Profiling result:

```

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:		98.10%	64.4583s	1	64.4583s	64.4583s	64.4583s	matMulGPU(unsigned int*, unsigned int*, unsigned int*, int*, int*, int*, int*, int, int, int, int)
		1.40%	921.59ms	1	921.59ms	921.59ms	921.59ms	[CUDA memcpy DtoH]
		0.47%	307.30ms	8	38.413ms	5.7910us	138.32ms	[CUDA memcpy HtoD]
		0.03%	18.419ms	1	18.419ms	18.419ms	18.419ms	[CUDA memset]
API calls:		98.89%	65.3804s	1	65.3804s	65.3804s	65.3804s	cudaMemcpy
		0.58%	386.23ms	2	193.11ms	6.6290us	386.22ms	cudaStreamCreate
		0.49%	326.70ms	8	40.837ms	20.126us	156.72ms	cudaMemcpyAsync
		0.02%	13.080ms	7	1.8686ms	25.425us	6.7470ms	cudaFree
		0.01%	7.6059ms	9	845.10us	5.7630us	4.4642ms	cudaMalloc
		0.00%	740.06us	2	370.03us	1.3170us	738.74us	cuDeviceGetPCIBusId
		0.00%	723.05us	2	361.52us	344.31us	378.74us	cuDeviceTotalMem
		0.00%	550.85us	202	2.7260us	126ns	123.41us	cuDeviceGetAttribute
		0.00%	67.440us	2	33.720us	8.1370us	59.303us	cudaDeviceSynchronize
		0.00%	59.749us	2	29.874us	23.872us	35.877us	cuDeviceGetName
		0.00%	41.699us	1	41.699us	41.699us	41.699us	cudaLaunchKernel
		0.00%	33.383us	1	33.383us	33.383us	33.383us	cudaMemset
		0.00%	22.023us	2	11.011us	4.0410us	17.982us	cudaStreamDestroy
		0.00%	1.8620us	4	465ns	135ns	1.0720us	cuDeviceGet
		0.00%	1.1170us	3	372ns	182ns	557ns	cuDeviceGetCount
		0.00%	496ns	2	248ns	223ns	273ns	cuDeviceGetUuid

```

real 2m20.456s
user 2m7.119s
sys 0m11.027s

```

```

Input reading done, n: 32768, m: 4, k-values: 33000000 and 33000000
Sorted input: 63.7081
==286== NVPROF is profiling process 286, command: ./exec input1 input2 outFile.bin
CSR converted
Time taken for matrix multiplication: 310.647
Writing output
Number of output non-zero blocks: 67108864
Writing done
==286== Profiling application: ./exec input1 input2 outFile.bin
==286== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	99.36%	306.862s	1	306.862s	306.862s	306.862s	matMulGPU(unsigned int*, unsigned int*, unsigned int*, int*, int*, int*, int*, int*, int, int, int, int)
	0.33%	1.02598s	8	128.25ms	10.687us	466.49ms	[CUDA memcpy HtoD]
	0.30%	931.87ms	1	931.87ms	931.87ms	931.87ms	[CUDA memcpy DtoH]
	0.01%	18.447ms	1	18.447ms	18.447ms	18.447ms	[CUDA memset]
API calls:	99.53%	307.794s	1	307.794s	307.794s	307.794s	cudaMemcpy
	0.34%	1.04535s	8	130.67ms	24.050us	484.88ms	cudaMemcpyAsync
	0.13%	392.57ms	2	196.28ms	3.5380us	392.57ms	cudaStreamCreate
	0.01%	16.332ms	7	2.3331ms	6.3520us	6.3657ms	cudaFree
	0.00%	11.147ms	9	1.2385ms	7.4930us	4.4978ms	cudaMalloc
	0.00%	1.0613ms	2	530.63us	1.6300us	1.0596ms	cuDeviceGetPCIBusId
	0.00%	814.36us	1	814.36us	814.36us	814.36us	cudaMemset
	0.00%	719.74us	2	359.87us	348.62us	371.12us	cuDeviceTotalMem
	0.00%	581.48us	202	2.8780us	128ns	130.50us	cuDeviceGetAttribute
	0.00%	86.239us	2	43.119us	11.294us	74.945us	cudaDeviceSynchronize
	0.00%	62.307us	2	31.153us	24.757us	37.550us	cuDeviceGetName
	0.00%	37.457us	2	18.728us	4.0820us	33.375us	cudaStreamDestroy
	0.00%	33.162us	1	33.162us	33.162us	33.162us	cudaLaunchKernel
	0.00%	2.1860us	4	546ns	150ns	1.2830us	cuDeviceGet
	0.00%	1.0410us	3	347ns	172ns	565ns	cuDeviceGetCount
	0.00%	498ns	2	249ns	207ns	291ns	cuDeviceGetUuid

```

real 7m7.282s
user 6m46.856s
sys 0m15.976s

```

### 5.3 Observation

Based on the above large outputs we can say that there is a massive speedup in GPU parallelism. The edge cases take 2min50secs(74 secs) and 7min7secs(310 secs) for  $m=8$  and  $4$  respectively. Also, the sparse matrix representation greatly improves the algorithm. Also based on different  $n$ ,  $m$ , and  $k$  values as in the graph, the scaling is almost proportional to the problem size. Though  $m$  as  $4$  or  $8$  does not have major effects when  $n$  is small.